

CIFAR-10 Neural Network Report

Task 1

Two transformations were created for the data: one for the train set and another for the test set.

Data augmentation and normalization were applied to the training set, only normalization was applied to the test set.

Augmentations [Train only]:

```
transforms.RandomHorizontalFlip() -> Default p=0.5  
transforms.RandomCrop(32, padding=4)  
transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2)  
transforms.RandomRotation(10)
```

Normalization [Train and Test]:

```
transforms.Normalize((0.4914, 0.4822, 0.4465), (0.247, 0.243, 0.261))
```

Normalization values specifically for the CIFAR-10 dataset were obtained from:
[GitHub Thread](#), [Code used to generate values](#)

Task 2

The model consists of two blocks followed by a classifier MLP. Regularization techniques were applied such as batch normalization and drop out were also used in between steps.

Block 1 [3 convolutional layers]:

```
Input X [batch_size,channels,height,width] -> [32,3,32,32]  
Take Average across dim=[2, 3]  
Pass through linear layer to obtain [a]  
Compute  $\text{Conv}_k(X)$  for each convolutional layer  
Compute  $O_1 = a_1 * \text{Conv}_1(X) + \dots + a_k * \text{Conv}_k(X)$   
Pass  $O_1$  through sequential
```

Block 2 [3 convolutional layers]:

```
Input  $O_1$  from Block 1 output [batch_size,channels,height/4,width/4] -> [32,256,8,8]  
Take Average across dim=[2, 3]  
Pass through MLP to obtain [a]  
Compute  $\text{Conv}_k(O_1)$  for each convolutional layer  
Compute  $O_2 = a_1 * \text{Conv}_1(O_1) + \dots + a_k * \text{Conv}_k(O_1)$   
Pass  $O_2$  through sequential
```

Classifier [MLP]:

Input O_2 from Block 2 output $[batch_size, channels, height/16, width/16] \rightarrow [32, 512, 2, 2]$

Take Average across $dim=[2, 3]$

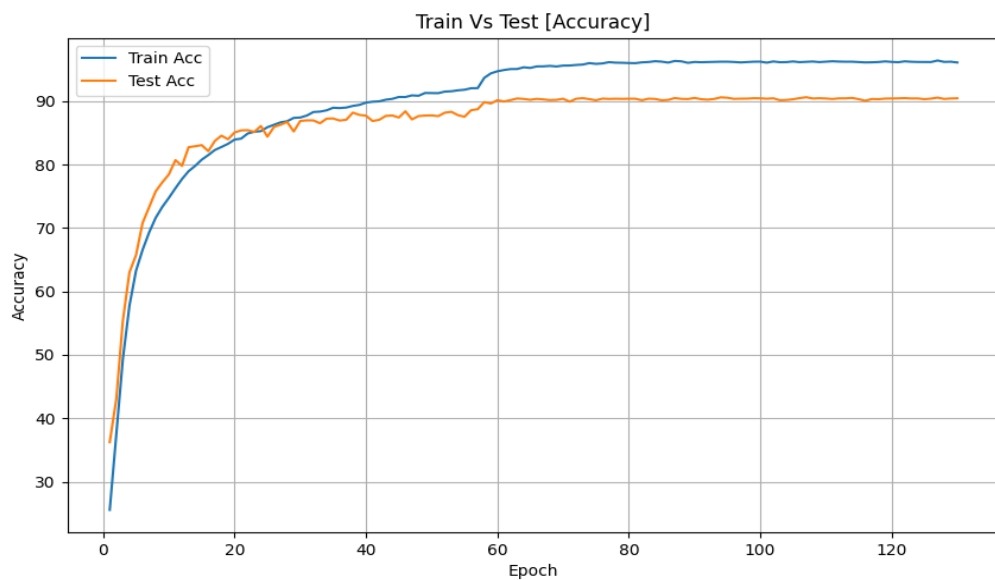
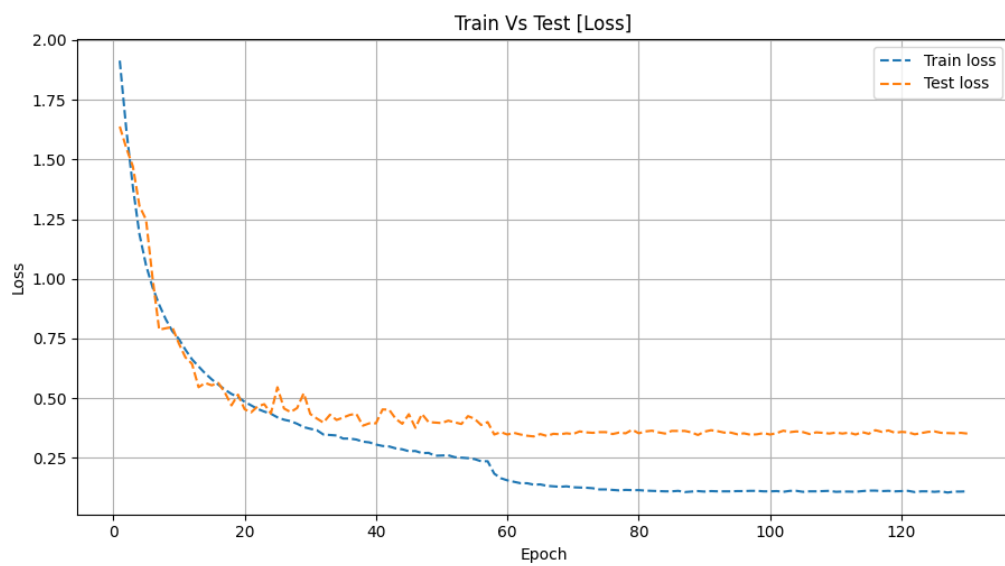
Pass through MLP to obtain 10 output features for classification

Task 3

`optimizer = torch.optim.Adam(net.parameters(), lr=0.01)`

`loss function = nn.CrossEntropyLoss()`

Task 4



Hyperparameters:

Batch Size -> 32

Epochs -> 130

Optimizer -> Adam

Learning Rate -> 0.01 (Added scheduler to reduce lr if test loss did not improve for 10 epochs)

Dropout Probability -> Convolutional layers (0.3), Linear layers (0.5)

Activation Function -> ReLU

Task 5

Model test accuracy peaked at epoch 107 of 130 with an accuracy of 90.63%.